



# DeepSeek-R1-Distill-Llama-70B Performance Analysis

## TP/PP Parallelism Configurations (Input: 8192 tokens, Output: 2048 tokens)



### Performance Scaling Efficiency & Cost Efficiency Analysis



### GPU Utilization Analysis

